



# Word Embeddings and Deep Learning for Spanish Twitter Sentiment Analysis

José Ochoa-Luna<sup>1</sup>(✉)  and Disraeli Ari<sup>2</sup>

<sup>1</sup> Department of Computer Science,  
Universidad Católica San Pablo, Arequipa, Peru  
jeochoa@ucsp.edu.pe

<sup>2</sup> Universidad Nacional de San Agustín, Arequipa, Peru  
darim@unsa.edu.pe

**Abstract.** Spanish is the third language most used on the internet. However, Natural Language Processing research in this language is still far below the level of other languages like English. The aim of this paper is to fill this gap in the literature and to provide a comprehensive assessment of Deep Learning applied to Spanish sentiment analysis. We focus on the polarity detection task which, in the context of Spanish Twitter messages, remains as a challenging task. To do so, we explore the combination of several Word representations (Word2Vec, Glove, Fastext) and Deep Neural Networks models. Unlike poor performance obtained by previous related work using Deep Learning for Spanish sentiment analysis, we show promising results. Our best setting combines three word embeddings representations, Convolutional Neural Networks and Recurrent Neural Networks. This setup allows us to obtain state-of-the-art results on the TASS/SEPLN 2017 Spanish Twitter benchmark dataset, in terms of accuracy and macro F1-measure.

**Keywords:** Spanish sentiment analysis · Deep learning · Word embeddings

## 1 Introduction

Online reviews are ubiquitous. On one hand, we have a wide variety of products and services being created every day. On the other hand, customers who have either purchased products or contracted services and, ultimately, left comments in social media. With the rapid growth of Twitter, Facebook, and online review sites, sentiment analysis draws growing attention from both research and industry communities [19].

Sentiment analysis, in its basic task called polarity detection, allows us to perform an automated analysis of millions of reviews and determine whether a given opinion is positive, negative or neutral. This area has been widely researched since 2002 [17]. In fact, it is one of the most active research areas in natural language processing, data mining and social media analytics [29].

Polarity detection has been addressed as a text classification problem thus, can be approached by supervised and unsupervised learning methods [31]. In the unsupervised approach, a vocabulary of positive and negative words is constructed so as to polarity is inferred according the similarity between vocabulary and opinionated words. The second approach is based on machine learning. Training data and labelled reviews are used to define a classifier [17]. This last approach relies heavily on feature engineering. However, recent learning representation paradigms perform this tasks automatically [16]. In this context, Machine Learning has recently become the dominant approach for sentiment analysis, due to availability of data, better models and hardware resources [30].

In this paper we adopt a Deep Learning approach for sentiment analysis. In particular, we aim at performing automated classification of short texts in sentences and Twitter messages for the Spanish language. This is challenging because of the limited contextual information that they normally contain.

In our proposed approach, sentence words are mapped to word representations. Three kinds of word representations (Word2vec [20], Glove [26], Fastext [4]) are used in our setting. This combination, which is novel for Spanish sentiment analysis, can be useful in several domains. Overall, our goal is to provide a general setup that can be applied with less effort in several contexts.

The Deep Learning architecture proposed is composed by a Convolutional Neural Network [15], a Recurrent Neural Network [13] and a final dense layer. In order to avoid overfitting, besides traditional dropout schemes, we propose a novel data augmentation approach. Data augmentation is useful for low resources languages such as Spanish—Specially for Spanish sentiment analysis.

Those design choices allow us to obtain state-of-the-art results, in terms of accuracy and macro F1 measure, on the InterTASS 2017 dataset. This dataset was proposed in the TASS workshop at SEPLN. In the last six years, this workshop has been the main source for Spanish sentiment analysis datasets and proposals [18].

The remainder of the paper is organized as follows. Related work is presented in Sect. 2. Background concepts are described in Sect. 3. Our proposal is presented in Sect. 4. Results are described in Sect. 5. Finally, Sect. 6 concludes the paper.

## 2 Related Work

There is a plethora of related works for sentiment analysis but, we are only interested in contributions for the Spanish language. Arguably, one of the most complete Spanish sentiment analysis systems was proposed by Brooke et al. [6], which had a linguistic approach. That approach integrated linguistic resources in a model to decide about polarity opinions [31]. Recent successful approaches for Spanish polarity classification have been mostly based on machine learning [10].

In the last six years, the TASS at SEPLN Workshop has been the main source for Spanish sentiment analysis datasets and proposals [11, 18]. Benchmarks for

both the polarity detection task and aspect-based sentiment analysis task have been proposed in several editions of this Workshop. Spanish Tweets have been emphasized.

Nowadays deep learning approaches emerge as powerful computational models that discover intricate semantic representations of texts automatically from data without feature engineering. These approaches have improved the state-of-the-art in many sentiment analysis tasks including sentiment classification of sentences/documents, sentiment extraction and sentiment lexicon learning [29]. However, these results have been mostly obtained for English Language. Since our proposal is based on Deep Learning, the related work that follows emphasizes these kinds of algorithms.

Arguably, the first approach using Deep Learning techniques for Spanish Sentiment Analysis was proposed in the TASS at SEPLN workshop in 2015 [32]. The authors presented one architecture that was composed by a RNN layer (LSTMs cells), a dense layer and a Sigmoid function as output. The performance over the general dataset was poor, 0.60 in terms of accuracy (the best result was 0.69 in TASS 2015).

The first Convolutional Neural Network approach for Spanish Sentiment Analysis was described by Segura-Bedmar et al. [28]. The CNN model proposed for sentiment analysis was mostly based on Kim's work [15]. It was comprised by only a single convolutional layer, followed by a max-pooling layer and a Softmax classifier as the final layer. Word embeddings were used in three ways: a learned word embedding from scratch and two pre-trained word2vec models. In terms of accuracy they obtained 0.64, which was far from the best result (0.72 was the best result in TASS 2016 [11]).

Another CNN approach for Spanish Sentiment Analysis was presented by Paredes et al. [25]. First, a preprocessing step (tokenization and normalization) was performed which was followed by a Word2vec embedding. Then, this model was comprised of a 2D convolutional layer, a max pooling and a final Softmax layer, i.e., it was also similar to Kim's work [15]. It was reported a F-measure of 0.887 over a non-public Twitter corpus of 10000 tweets.

Most of the Deep Learning approaches for Spanish sentiment analysis have been presented in TASS 2017 [18]. For instance, Rosa et al. [27] used word embeddings within two approaches: SVM (with manually crafted features) and Convolutional Neural Networks. Pre-trained Word2vec, Glove and Fastext embeddings were used. Unlike our approach, these embeddings were used separately. In fact, the best results of this paper were obtained using Word2vec. When CNN was employed, unidimensional convolutions were performed. While several convolutional layers were tested, the best model had three convolutional layers, using 2, 3 and 4 word filters. However, their best results were obtained when combined SVM with CNN. They simply used a decision rule based on both probability results. Interesting results, in terms of accuracy, were obtained. It was reported a 0.596 value for the InterTASS dataset (the best accuracy result was 0.608 for TASS 2017 [18]).

Garcia-Vega et al. [12] used word embeddings with shallow classifiers. In addition, they also tested recurrent neural networks with LSTM nodes and a dense layer. Two kinds of experiments were performed using word embeddings and TFIDF values as inputs. Both experiments obtained poor results (0.333 and 0.404 in terms of accuracy for the 2017 InterTASS dataset).

Araque et al. [1] explored recurrent neural networks in two ways (i) a set of LSTM cells whose input were word embeddings, (ii) a combination of input word vector and polarity values obtained from a sentiment lexicon. As usual, a last dense layer with a Softmax function was used as final output. While interesting, experimental results showed that the best performance was obtained by the second model, i.e., LSTM + Lexicon + dense layer. They obtained 0.562 for 2017 InterTASS dataset, in terms of accuracy. This value was far from the top results.

In the last years, the best results were obtained for the ELiRF group [14]. In TASS 2017, they obtained the second best result for the InterTSS task, 0.607, in terms of accuracy (The first place presented an ensemble approach [7]). It is worth noting that ELiRF best results were obtained using a Multilayer perceptron (MLP) with word embeddings as inputs. This MLP had two layers with ReLu activation functions. A Second approach used a stack of CNN and LSTM models with pre-trained word embeddings. The architecture was composed by one convolutional layer, 64 LSTM cells and a fully connected MLP layer. This last architecture had a poor performance (0.436 in terms of Accuracy).

### 3 Background

#### 3.1 Sentiment Analysis

Sentiment analysis (also known as opinion mining) is an active research area in natural language processing [30]. Sentiment classification is a fundamental and extensively studied area in sentiment analysis. It targets at determining the sentiment polarity (positive or negative) of a sentence (or a document) based on its textual content [29]. Polarity classification tasks have usually based on two main approaches [5]: a supervised approach, which applies machine learning algorithms in order to train a polarity classifier using a labelled corpus; an unsupervised approach, semantic lexicon-based, which integrates linguistic resources in a model in order to identify the polarity of the opinions.

Since the performance of a machine learner heavily depends on the choices of data representation, many studies devote to building powerful feature extractor with domain expert and careful engineering [22].

As stated by Liu [17], sentiment analysis has been researched at three levels:

- Document level: The task at this level is to classify whether a whole opinion document expresses a positive or negative sentiment [24]
- Sentence level: The task at this level goes to the sentences and determines whether each sentence expressed a positive, negative, or neutral opinion. Neutral usually means no opinion.

- Entity and Aspect level [3]: Both the document level and the sentence level analyses do not discover what exactly people liked and did not like. Aspect level performs finer-grained analysis.

### 3.2 Deep Neural Networks

Several deep neural network approaches have been successfully applied to sentiment analysis in the last years [33]. However, these results have been mostly obtained for English Language [18]. In this section we only focus on word representations, Convolutional Neural Networks (CNNs) and Recurrent Neural Networks (RNNs). They are the main building blocks of our proposal.

**Word Representations (Word2vec, Glove, Fastext).** Nowadays, word representations are paramount for sentiment analysis [33]. In order to model text words as features within a machine learning framework, a common approach is to encode words as discrete atomic symbols. These encodings are arbitrary and provide no useful information to the system regarding the relationships that may exist between the individual symbols [30].

The discrete representation has some problems such as missing new words. This representation also requires human labor to create and adapt. It is also hard to compute accurate word similarity and is quite subjective. To cope with these problems, the distributional similarity based representations propose to represent a word by means of its neighbors, its context [29].

Word2vec [20] is a particularly computationally-efficient predictive model for learning word embeddings from raw text. It takes a vector with several hundred dimensions where each word is represented by a distribution of weights across those elements [2, 8]. Thus, instead of a one-to-one mapping between an element in the vector and a word, the representation of a word is spread across all the elements in the vector. In addition, each element in the vector contributes to the definition of many words. Such a vector comes to represent in some abstract way the “meaning” of a word. And simply by examining a large corpus it is possible to learn word vectors that are able to capture the relationships between words in a surprisingly expressive way.

Unlike Word2vec, Glove [26] seeks to make explicit what Word2vec does implicitly: encoding meaning as vector offsets in an embedding space. In Glove, it is stated that the ratio of the co-occurrence probabilities of two words (rather than their co-occurrence probabilities themselves) is what contains information and so look to encode this information as vector differences.

Instead of directly learning a vector representation for a word, Fastext [4] learns a representation for each character n-gram. In this sense, each word is represented as a bag of characters n-grams. Thus, the overall word embedding is a sum of these characters n-grams. The advantage of Fastext is that generates better embeddings for rare and out-of-corpus words. By using different n-grams, Fastext explores key structural components of words.

**Convolutional Neural Networks.** While Convolutional Neural Networks (CNN) have been primarily applied to image processing, they have also been used for NLP tasks [15].

In the image context [16], given a raw input (2D arrays of pixel intensities), several *convolutional* layers allow us to capture features images at several abstraction levels. In this context, a discrete convolution takes a filter matrix and multiply its values element-wise with the original matrix, then sum them up. To get the full convolution we do this for each element by sliding the filter over the whole matrix.

The convolved map feature denotes a level of abstraction obtained after the convolution operations (there are also ReLU activation, Pooling and Softmax layers). CNN exploits the property that many natural signals are compositional hierarchies: higher-level features are obtained by composing lower-level ones. In images, local combinations of edges form motifs, motifs assemble into parts, and parts from objects [16]. All this learning representation is performed in an unsupervised manner. The amount of filters and convolutional layers denote how rich features and abstraction levels we wish to obtain from images.

Conversely, if we wish to apply CNNs in natural language tasks several changes are needed [15]. Text data is tokenized and must be encoded as numbers (input numerical variables are usual in neural networks algorithms). In the last five years, word embeddings representations (but also character and paragraph) have been preferred. This is due to semantical/syntactical similarity is better expressed in a distributed manner [20].

A sentence can be represented as a matrix. The sentence length denotes the number of rows and the word embedding dimension denotes the number of columns. This allows us to perform discrete convolutions as in the image case (2D input matrix). However, one must be careful when defining filter sizes which usually have the same width as word embeddings [15].

Instead of working with 2D representation, we may also work with 1D representation, i.e., to concatenate several word embeddings in a long vector and then apply several convolutional layers.

**Recurrent Neural Networks.** Recurrent Neural Networks (RNN) [9] are a kind of neural network that makes it possible to model long-distance dependencies among variables. Therefore, RNNs are best suited for tasks that involve sequential inputs, such as speech and language [16]. RNNs process an input sequence one element at a time, maintaining in their hidden units a state vector that implicitly contains information about the history of all the past elements of the sequence. To do so, a connection is added that references the previously hidden states  $h_{t-1}$  when computing hidden state  $h$ , formally [23]:

$$h_t = \tanh(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$h_t = 0$  when the initial step is  $t = 0$ . The only difference from the hidden layer in a standard neural network is the addition of the connection  $W_{hh}h_{t-1}$  from

the hidden state at time step  $t - 1$  connecting to that at time step  $t$ . Since this is a recursive equation that uses  $h_{t-1}$  from the previous time step.

In the context of Sentiment Analysis, an opinionated sentence is a sequence of words. Thus, RNNs are suitable for modeling this input sequence [13]. Similar to CNNs, the input is given as words (character) embeddings which can be learned during training or may also be pre-trained (Glove, Word2vec, Fastext).

Each word is mapped to a word embedding which is the input at every time step of the RNN. The maximum sequence length denotes the length of the recurrent neural network. Each hidden state models the dependence among a current word and all the precedent words. Usually the final hidden state, which ideally denotes all the encoded sentence, is connected to a dense layer so as to perform sentiment classification [13].

RNNs are very powerful dynamic systems, but training them has proved to be problematic. The backpropagated gradients either grow or shrink at each time step. Thus, over many time steps they typical explode or vanish. A sequence of words comprise a sequence of RNNs cells. These cells can have some gate mechanism in order to avoid gradient vanishing longer sequences. In this setting Long Short Term Memory Cells (LSTM) or Gated Recurrent Units (GRU) are common choices [23].


## 4 Proposal

The aim of this paper is to explore several Deep Learning algorithms possibilities in order to perform sentiment analysis. The focus is to tackle the polarity detection task for Spanish Tweets. In this sense, some models were tested. Details of these experiments are given in Sect. 5.

In this section, we present our best pipeline for Spanish sentiment analysis given short texts. Basically, it is composed by Word embeddings, CNN and RNN models. The pipeline is showed in Figure 1. A concise description is given as follows.

- Basic pre-processing is performed as the focus is given to data augmentation.
- The input is a sequence of words—a short opinionated sentence. These words are mapped to three pre-trained Spanish word embeddings (Word2vec, Glove, Fastext).
- The three channels are the input to a 3D Convolutional Neural Network. After several convolutional and max pooling layers we obtain a feature vector of a given length.
- The feature vector obtained from the CNN is mapped to a sequence and passed to a RNN. It is a simple RNN model, with LSTM cells.
- The final hidden state of the RNN is completely connected to a dense layer to train a classifier.

Further details about these design choices are given as follows.



**Fig. 1.** Pipeline of our proposal: Word Embeddings+CNN+RNN.

#### 4.1 Data Augmentation

In general a few pre-processing steps are performed over raw data. Since we have few training examples in Spanish and Deep Learning techniques are susceptible to overfitting, we would rather focus on data augmentation. We propose a novel approach for data augmentation. Basically, we identify nouns, adjectives and verbs on sentences by performing Part-Of-Speech tagging<sup>1</sup>. By doing so, we emphasize tokens that are prone to be opinionated words. Then, more examples are created by combining bigrams and trigrams from the former tokens. In addition, we augment data based on word synonyms [33]. Opinionated words are replaced by synonyms. Overall, this process allowed us to obtain better generalization results.

#### 4.2 Word Embeddings Choice

One of the main contributions of this paper was to find the best word embedding setting. We have trained Word2vec and Glove embeddings on Spanish corpus. Moreover, we have used a pre-trained Fasttext embedding. At the end, empirical tests allowed us to decide for using these three mappings as channels in our CNN building block. In the context of Spanish sentiment analysis, none of the previous works had used three embedding channels in CNNs before. We argue that by using Word2vec, Glove and Fasttext at the same time, we were able to take advantage of several “word meanings”. These word meanings are accordingly linked to word context, co-occurrence probabilities and rare words representations.

<sup>1</sup> The following tool was used to perform POS tagging: <http://www.cis.uni-muenchen.de/~schmid/tools/TreeTagger/>.



### 4.3 CNN Architecture

Our CNN architecture is based on Kim's work [15]. Since three word embeddings are used, the first convolutional layer receives a 3D input. Each word embedding has 300 dimensions. Filters have the same width as embeddings dimension and we perform convolutions from 1 to 5 words. The number of filters is 300. The pooling layer allows us to control the desired feature vector obtained.

Some other hyper paremeters are given as follows:

- batch size: 32
- dropout keep probability: 0.5
- filter sizes: 3,4,5
- hidden unit: 300
- maximum pooling size: 4
- number of epochs: 10

### 4.4 RNN Architecture

The RNN receives a CNN vector as input and LSTMs cells are defined accordingly. There is one layer of LSTM cells, these cells have 300 hidden units. The last hidden state is fully connected to a dense layer which allows us to define a classifier [13].

## 5 Experiments

Experiments were performed using Deep Learning algorithms. CNNs and RNNs were tested separately. Our best result was obtained by composing word embeddings, CNNs and RNNs. We first describe the benchmark dataset used. Then, accuracy results are showed.

### 5.1 Dataset

The dataset used to perform comparisons was InterTASS which is a collection of Tweets written in Spanish. It is composed of more than 3,000 tweets annotated at four opinion intensity level (positive, neutral, negative and none). We only focus on task 1. This task aims to evaluate polarity classification at tweet level. This shared task has been proposed at the TASS SEPLN workshop in 2017 [18]. We have used this dataset since it is the most recent benchmark that allows us to compare among Deep Learning approaches for Spanish sentiment analysis. The dataset is further detailed in Table 1.

**Table 1.** InterTASS dataset (TASS 2017)

Corpus	Tweets
Training	1,008
Development	506
Test	1,899
Total	3,413

## 5.2 Results

We have implemented several deep neural networks models, and the dataset InterTASS 2017 was used for training. For this implementation we use Tensorflow<sup>2</sup>. In order to find the best hyper parameters, we have used a ten-fold cross validation process. The test set has only been used to report results. In Table 2 we report results in terms of accuracy.

A first attempt was to test several RNNs models (many-to-one architecture , single layer, multilayer, bidirectional). The reported model, RNN in Table 2, has a many-to-one architecture. The input is a sequence of words and the output is the resulting polarity. There is only a hidden layer, and the input is a pre-trained sequence of Word2vec embeddings.

A second attempt was to test several CNN models, i.e., 1D CNN, 2D CNN and 3D CNNs, until 4 convolutional/pooling layers. The reported model, CNN in Table 2, is a 3D CNN. Thus, the input received three channels of pre-trained word embeddings. It had only three layers: a convolutional, a pooling and a dense layer.

It is worth noting that our best result was obtained by the model described in Sect. 3 (CNN+RNN in Table 2). This is a combination of a 3D CNN and a many-to-one RNN. A 3D CNN architecture whose outputs where mapped to a sequence of LSTM cells. Our data augmentation scheme was also used in order to avoid overfitting.

**Table 2.** Deep learning approaches results on InterTASS dataset (TASS 2017)

Our DL attempts	Accuracy
<b>CNN+RNN</b>	<b>0.609</b>
CNN	0.5552
RNN	0.4972

In Table 3, we compare, in terms of accuracy and Macro F1 measure, our best model with the state-of-the-art InterTASS 2017 results. It is worth noting that

<sup>2</sup> <https://www.tensorflow.org/>.

**Table 3.** State-of-the-art results on InterTASS dataset (TASS 2017)

System	Accuracy	Macro F1
<b>CNN+RNN (our approach)</b>	<b>0.609</b>	<b>0.551</b>
jacerong-run1 [7]	0.608	0.459
ELIRF-UPV-run1 [14]	0.607	0.493
RETUYT-svm cnn [27]	0.596	0.471
tecnolenguasent [21]	0.595	0.456

our approach outperforms all the other approaches. In addition, our proposal is the only top result using a Deep Learning approach.

## 6 Conclusion

Despite being one of the three most used languages at Internet, Spanish has had few resources developed for natural language processing tasks. Unlike English sentiment analysis, Deep Learning approaches were unable to obtain state-of-the-art results on Spanish benchmark datasets in the past. The aim of this work was to provide empirical evidence on the performance of Deep Learning algorithms for Spanish Twitter sentiment analysis. Thus, we have showed that a combination of data augmentation, at least three kinds of word embeddings, a 3D Convolutional Neural Network, followed by a Recurrent Neural Network allows us to obtain state-of-the-art results, in terms of accuracy and macro F1 measure on the InterTASS 2017 benchmark. In addition, this setup can be easily adapted to other domains.

## References

1. Araque, O., Barbado, R., Sanchez-Rada, J.F., Iglesias, C.A.: Applying recurrent neural networks to sentiment analysis of Spanish tweets. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 71–76 (2017)
2. Bengio, Y., Ducharme, R., Vincent, P., Janvin, C.: A neural probabilistic language model. *J. Mach. Learn. Res.* **3**, 1137–1155 (2003). <http://dl.acm.org/citation.cfm?id=944919.944966>
3. Blair-goldensohn, S., Neylon, T., Hannan, K., Reis, G.A., Mcdonald, R., Reynar, J.: Building a sentiment summarizer for local service reviews. In: *NLP in the Information Explosion Era* (2008)
4. Bojanowski, P., Grave, E., Joulin, A., Mikolov, T.: Enriching word vectors with subword information. *Trans. Assoc. Comput. Linguist.* **5**, 135–146 (2017)
5. Brody, S., Elhadad, N.: An unsupervised aspect-sentiment model for online reviews. In: *Human Language Technologies: The 2010 Annual Conference of the North American Chapter of the Association for Computational Linguistics, HLT 2010*, pp. 804–812. Association for Computational Linguistics, Stroudsburg (2010). <http://dl.acm.org/citation.cfm?id=1857999.1858121>

6. Brooke, J., Tofiloski, M., Taboada, M.: Cross-linguistic sentiment analysis: from English to Spanish. In: Proceedings of RANLP 2009, pp. 50–54 (2009)
7. Ceron-Guzman, J.A.: Classifier ensembles that push the state-of-the-art in sentiment analysis of Spanish tweets. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 59–64 (2017)
8. Collobert, R., Weston, J., Bottou, L., Karlen, M., Kavukcuoglu, K., Kuksa, P.P.: Natural language processing (almost) from scratch. CoRR abs/1103.0398 (2011). <http://arxiv.org/abs/1103.0398>
9. Elman, J.L.: Finding structure in time. Cogn. Sci. **14**(2), 179–211 (1990)
10. Garcia, M., Martinez, E., Villena, J., Garcia, J.: TASS 2015 - the evolution of the spanish opinion mining systems. Procesamiento de Lenguaje Natural **56**, 33–40 (2016)
11. Garcia-Cumbreras, M.A., Villena-Roman, J., Martinez-Camara, E., Diaz-Galiano, M., Martin-Valdivia, T., Ureña Lopez, A.: Overview of TASS 2016. In: Proceedings of TASS 2016: Workshop on Sentiment Analysis at SEPLN, pp. 13–21 (2016)
12. Garcia-Vega, M., Montejo-Raez, A., Diaz-Galiano, M.C., Jimenez-Zafra, S.M.: Sinai in TASS 2017: tweet polarity classification integrating user information. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 91–96 (2017)
13. Graves, A.: Supervised Sequence Labelling with Recurrent Neural Networks. Studies in Computational Intelligence, vol. 385. Springer, Berlin (2012). <https://doi.org/10.1007/978-3-642-24797-2>. <https://cds.cern.ch/record/1503877>
14. Hurtado, L.F., Pla, F., Gonzalez, J.A.: ELiRF-UPV at TASS 2017: sentiment analysis in twitter based on deep learning. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 29–34 (2017)
15. Kim, Y.: Convolutional neural networks for sentence classification. In: Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing, EMNLP 2014, Doha, Qatar, 25–29 October 2014, A meeting of SIGDAT, a Special Interest Group of the ACL, pp. 1746–1751 (2014). <http://aclweb.org/anthology/D/D14/D14-1181.pdf>
16. LeCun, Y., Bengio, Y., Hinton, G.: Deep learning. Nature **521**(7553), 436–444 (2015)
17. Liu, B.: Sentiment Analysis and Opinion Mining. Morgan and Claypool Publishers (2012)
18. Martinez-Camara, E., Diaz-Galiano, M., Garcia-Cumbreras, M.A., Garcia-Vega, M., Villena-Roman, J.: Overview of TASS 2017. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 13–21 (2017)
19. McGlohon, M., Glance, N., Reiter, Z.: Star quality: Aggregating reviews to rank products and merchants. In: Proceedings of Fourth International Conference on Weblogs and Social Media (ICWSM) (2010)
20. Mikolov, T., Sutskever, I., Chen, K., Corrado, G.S., Dean, J.: Distributed representations of words and phrases and their compositionality. In: Burges, C.J.C., Bottou, L., Welling, M., Ghahramani, Z., Weinberger, K.Q. (eds.) Advances in Neural Information Processing Systems 26, pp. 3111–3119. Curran Associates, Inc. (2013). <http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf>
21. Moreno-Ortiz, A., Perez-Hernandez, C.: Tecnolengua Lingmotif at TASS 2017: Spanish twitter dataset classification combining wide-coverage lexical resources and text features. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 35–42 (2017)

22. Narayanan, V., Arora, I., Bhatia, A.: Fast and accurate sentiment classification using an enhanced naive bayes model. In: Yin, H., et al. (eds.) IDEAL 2013. LNCS, vol. 8206, pp. 194–201. Springer, Heidelberg (2013). [https://doi.org/10.1007/978-3-642-41278-3\\_24](https://doi.org/10.1007/978-3-642-41278-3_24)
23. Neubig, G.: Neural machine translation and sequence-to-sequence models: A tutorial. CoRR abs/1703.01619 (2017). <http://arxiv.org/abs/1703.01619>
24. Pang, B., Lee, L.: Opinion mining and sentiment analysis. Found. Trends Inf. Retr. **2**(1–2), 1–135 (2008). <https://doi.org/10.1561/15000000011>
25. Paredes-Valverde, M.A., Colomo-Palacios, R., Salas-Zarate, M.D.P., Valencia-Garcia, R.: Sentiment analysis in Spanish for improvement of products and services: a deep learning approach. Sci. Program. **6**, 1–6 (2017)
26. Pennington, J., Socher, R., Manning, C.D.: Glove: global vectors for word representation. In: Empirical Methods in Natural Language Processing (EMNLP), pp. 1532–1543 (2014). <http://www.aclweb.org/anthology/D14-1162>
27. Rosa, A., Chiruzzo, L., Etcheverry, M., Castro, S.: RETUYT in TASS 2017: sentiment analysis for Spanish tweets using SVM and CNN. In: Proceedings of TASS 2017: Workshop on Sentiment Analysis at SEPLN, pp. 77–83 (2017)
28. Segura-Bedmar, I., Quiros, A., Martínez, P.: Exploring convolutional neural networks for sentiment analysis of Spanish tweets. In: Proceedings of the 15th Conference of the European Chapter of the Association for Computational Linguistics: vol. 1, Long Papers, pp. 1014–1022. Association for Computational Linguistics (2017). <http://aclweb.org/anthology/E17-1095>
29. Tang, D., Wei, F., Qin, B., Yang, N., Liu, T., Zhou, M.: Sentiment embeddings with applications to sentiment analysis. IEEE Trans. Knowl. Data Eng. **28**(2), 496–509 (2016)
30. Tang, D., Qin, B., Liu, T.: Deep learning for sentiment analysis: successful approaches and future challenges. Wiley Interdisc. Rev.: Data Mining Knowl. Discov. **5**(6), 292–303 (2015)
31. Turney, P.D.: Thumbs up or thumbs down?: semantic orientation applied to unsupervised classification of reviews. In: Proceedings of the 40th Annual Meeting on Association for Computational Linguistics, ACL 2002, pp. 417–424. Association for Computational Linguistics, Stroudsburg (2002). <https://doi.org/10.3115/1073083.1073153>
32. Vilares, D., Doval, Y., Alonso, M.A., Gomez-Rodriguez, C.: LyS at TASS 2015: deep learning experiments for sentiment analysis on Spanish tweets. In: Proceedings of TASS 2015: Workshop on Sentiment Analysis at SEPLN, pp. 47–52 (2015)
33. Zhang, L., Wang, S., Liu, B.: Deep learning for sentiment analysis: a survey. CoRR abs/1801.07883 (2018). <http://arxiv.org/abs/1801.07883>